


# Word Embeddings based on Neural Networks

Word Embeddings represent words  
as **low-dimensional vectors**  
in mathematical space and  
**capture** their semantic and  
syntactic **meaning**.

# Natural Language Processing


## Word Embedding Approaches



# Natural Language Processing

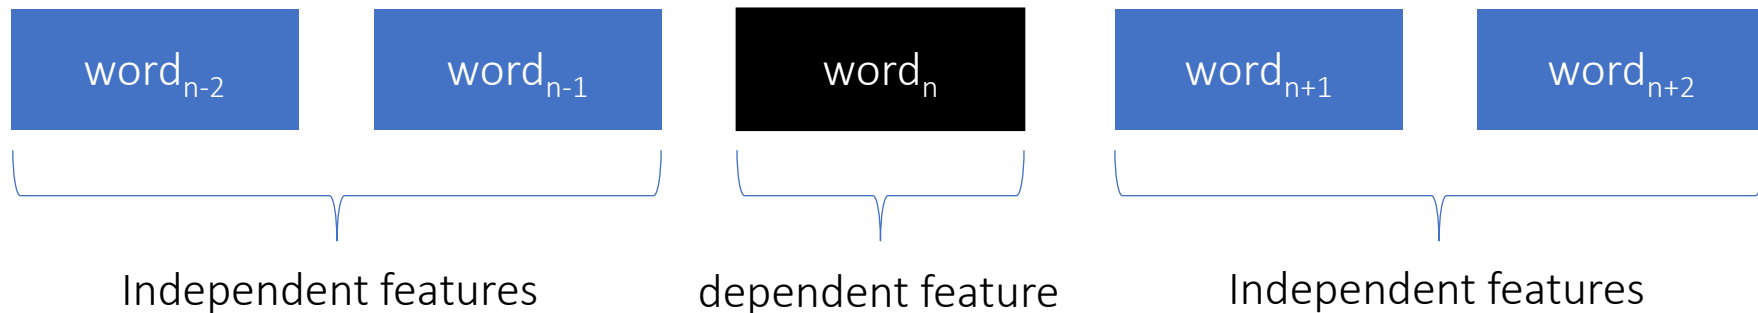
## Neural Network based Embeddings

- Aim to
  - Capture context / meaning
  - Capture similarity to other words
  - Reduce dimension
  - Avoid memory issues
- Developed based on Neural Networks



# Natural Language Processing

Word2Vec: Continuous Bag of Words

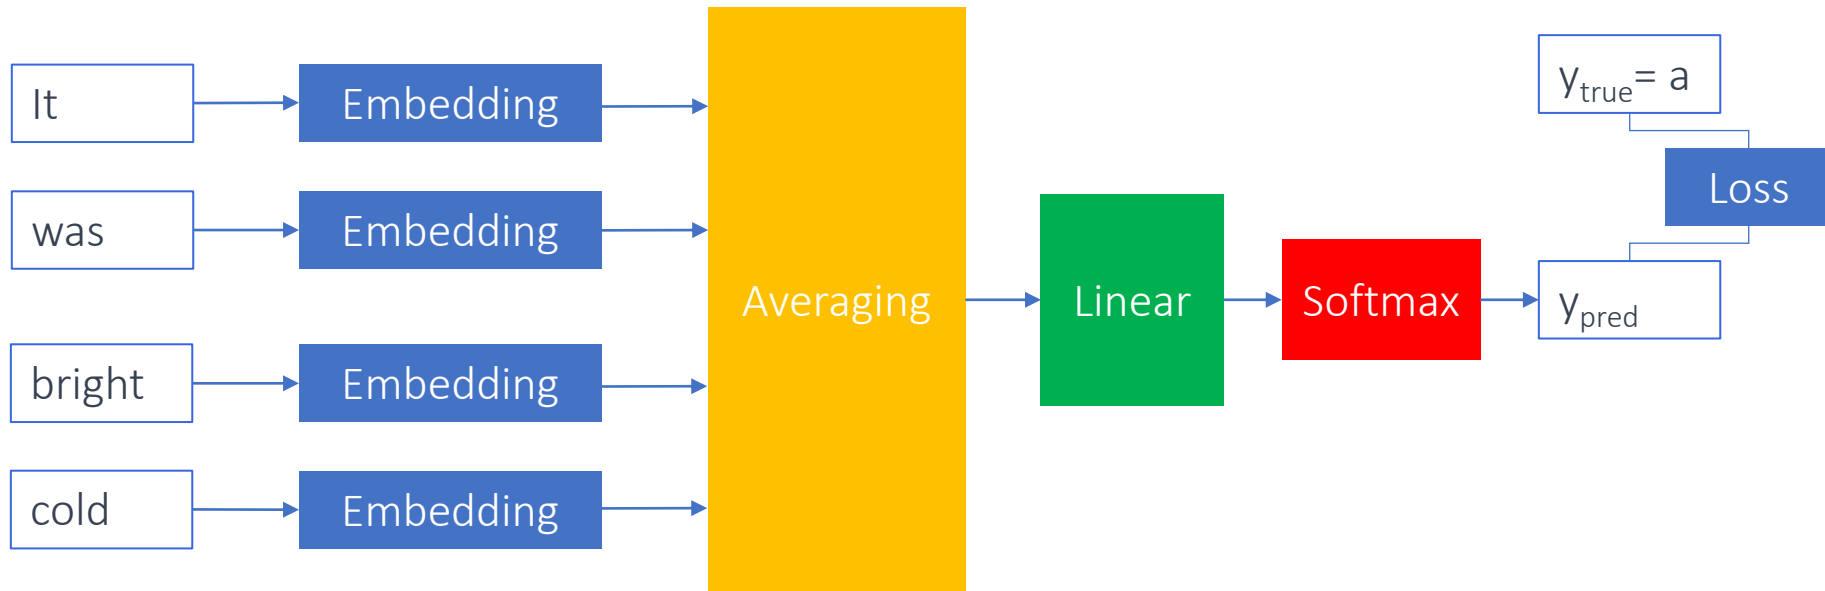


Independent Features	Dependent Feature“
[“It”, “was”, “bright”, “cold”]	“a”
[“was”, “a”, “cold”, “day”]	„bright“
...	

# Natural Language Processing

Word2Vec: Continuous Bag of Words Model

Inputs



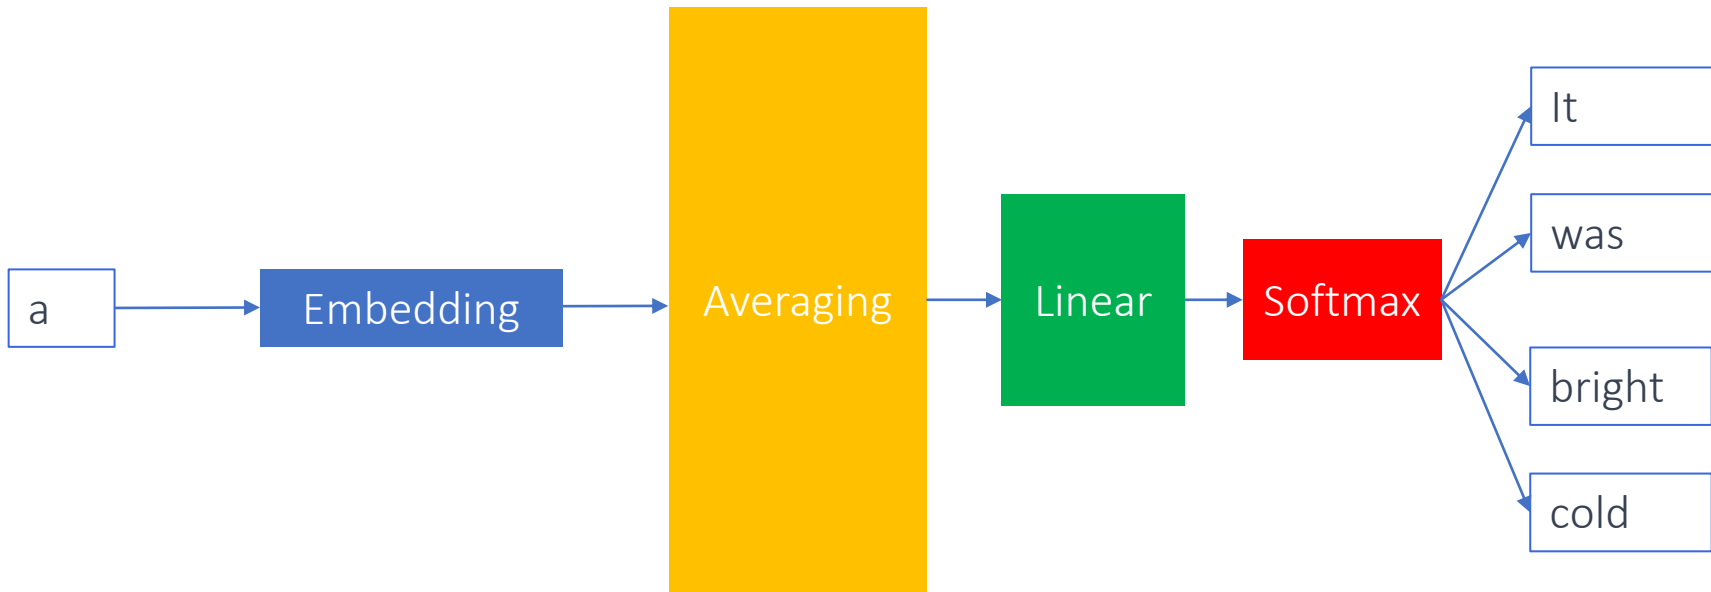
# Natural Language Processing

Word2Vec: Skip Gram

Inputs


Model

Outputs



# Natural Language Processing

## Word Embedding Approaches






# Natural Language Processing

## GloVe


- Global Vectors for Word Representations
- Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. [GloVe: Global Vectors for Word Representation](#)
- based on co-occurrence matrix of words in a corpus, which counts how often words appear together in the same context.
- constructs a matrix of word co-occurrence counts and then factorizes this matrix to obtain word embeddings
- factorization based on singular value decomposition (SVD)
- resulting embeddings are dense, low-dimensional vectors
- Encode words as vector of other words



Source: <https://nlp.stanford.edu/projects/glove/>

# Natural Language Processing

## Word Embedding Approaches



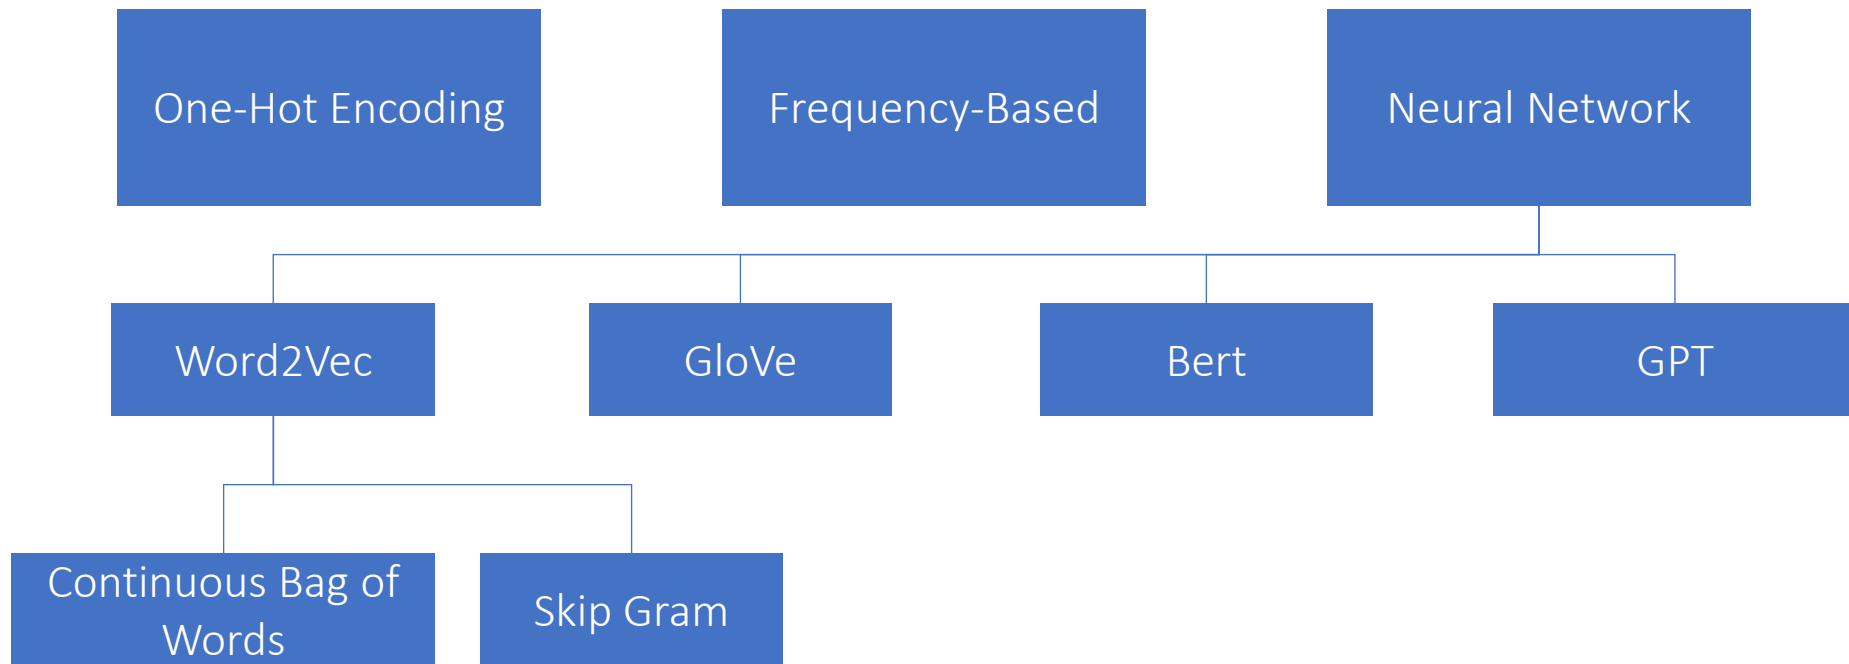
# Natural Language Processing

## BERT

- **Bidirectional Encoder Representations from Transformers**
- Developed by Google in 2018
- Pre-trained word embedding
- Based on Transformers
- Applies „masked language modeling“ – masking some words in sentence and learn to predict them
- Applies „next sentence prediction“ – model predicts whether two sentences are similar in a text
  
- Original variants: BERT-base (110m parameters, 440MB) and BERT-large (340m parameters, 1.3GB)
- Other variants: RoBERTa, ALBERT, ELECTRA, ...

# Natural Language Processing

## Word Embedding Approaches



# Natural Language Processing

GPT

- **Generative Pre-trained Transformers**
- Developed by OpenAI
- Not strictly a word embedding, but contextualized word embedding
- Unique embedding for each occurrence of a word based on surrounding words in text
- Applies Transformer architecture
- GPT-3 has 175 billion parameters

