

PyTorch Lightning

PyTorch Lightning

Introduction

- Python package
- framework built on top of PyTorch
- claims „making coding complex networks simple“
- nice features
 - logging of training / validation metrics
 - creating checkpoints
 - early stopping
 - training on multiple GPUs, TPUs, and CPUs
- Installation:

```
pip install pytorch-lightning
```



PyTorch Lightning

Introduction

- requires a class with four functions:
 - `__init__()`
 - `forward()`
 - `configure_optimizers()`
 - `training_step()`
- optional functions
 - `prepare_data()`
 - `validation_step()`
 - `test_step()`
 - `predict_step()`



PyTorch Lightning

Comparison PyTorch vs. PyTorch Lightning

your package imports

your package imports

```
import pytorch_lightning as pl
```

PyTorch

PyTorch Lightning

PyTorch Lightning

Comparison PyTorch vs. PyTorch Lightning

Class Inheritance

```
class LinearRegressionTorch(nn.Module):
    def __init__(self, input_size, output_size):
        super(LinearRegressionTorch, self).__init__()
        self.linear = nn.Linear(input_size, output_size)

    def forward(self, x):
        return self.linear(x)

model = LinearRegressionTorch(input_size=1, output_size=1)
model.train()

Zelle ausführen | Oben ausführen | Zelle debuggen
# %% Mean Squared Error
loss_fun = nn.MSELoss()

Zelle ausführen | Oben ausführen | Zelle debuggen
# %% Optimizer
learning_rate = 0.02
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

Zelle ausführen | Oben ausführen | Zelle debuggen
# %% perform training
losses = []
slope, bias = [], []
number_epochs = 1000
for epoch in range(number_epochs):
    for j, (X, y) in enumerate(train_loader):
        # optimization
        optimizer.zero_grad()
        # forward pass
        y_pred = model(X)
        # compute loss
        loss = loss_fun(y_pred, y)
        losses.append(loss.item())
        loss.backward()
        # update weights
        optimizer.step()
    # store loss
    losses.append(float(loss.data))
```

PyTorch

```
class LitlinearRegression(pl.LightningModule):
    def __init__(self, input_size, output_size):
        super(LitlinearRegression, self).__init__()
        self.linear = nn.Linear(input_size, output_size)
        self.loss_fun = nn.MSELoss()

    def forward(self, x):
        return self.linear(x)

    def configure_optimizers(self):
        learning_rate = 0.02
        optimizer = torch.optim.SGD(self.parameters(), lr=learning_rate)
        return optimizer

    def training_step(self, train_batch, batch_idx):
        X, y = train_batch

        # forward pass
        y_pred = model(X)

        # compute loss
        loss = self.loss_fun(y_pred, y)
        self.log('train_loss', loss, prog_bar=True)

        return loss
```

```
Zelle ausführen | Oben ausführen | Zelle debuggen
# %% model instance and training
# model instance
model = LitlinearRegression(input_size=1, output_size=1)

# training
trainer = pl.Trainer(gpus=1, precision=16, max_epochs=100)
trainer.fit(model, train_loader)
```

PyTorch Lightning

PyTorch Lightning

Comparison PyTorch vs. PyTorch Lightning

Optimizer

```
class LinearRegressionTorch(nn.Module):
    def __init__(self, input_size, output_size):
        super(LinearRegressionTorch, self).__init__()
        self.linear = nn.Linear(input_size, output_size)

    def forward(self, x):
        return self.linear(x)

model = LinearRegressionTorch(input_size=1, output_size=1)
model.train()

Zelle ausführen | Oben ausführen | Zelle debuggen
# %% Mean Squared Error
loss_fun = nn.MSELoss()

Zelle ausführen | Oben ausführen | Zelle debuggen
# %% Optimizer
learning_rate = 0.02
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

Zelle ausführen | Oben ausführen | Zelle debuggen
# %% perform training
losses = []
slope, bias = [], []
number_epochs = 1000
for epoch in range(number_epochs):
    for j, (X, y) in enumerate(train_loader):
        # optimization
        optimizer.zero_grad()
        # forward pass
        y_pred = model(X)
        # compute loss
        loss = loss_fun(y_pred, y)
        losses.append(loss.item())
        loss.backward()
        # update weights
        optimizer.step()
    # store loss
    losses.append(float(loss.data))
```

PyTorch

```
class LitLinearRegression(pl.LightningModule):
    def __init__(self, input_size, output_size):
        super(LitLinearRegression, self).__init__()
        self.linear = nn.Linear(input_size, output_size)
        self.loss_fun = nn.MSELoss()

    def forward(self, x):
        return self.linear(x)

    def configure_optimizers(self):
        learning_rate = 0.02
        optimizer = torch.optim.SGD(self.parameters(), lr=learning_rate)
        return optimizer

    def training_step(self, train_batch, batch_idx):
        X, y = train_batch

        # forward pass
        y_pred = model(X)

        # compute loss
        loss = self.loss_fun(y_pred, y)
        self.log('train_loss', loss, prog_bar=True)
        return loss

Zelle ausführen | Oben ausführen | Zelle debuggen
# %% model instance and training
# model instance
model = LitLinearRegression(input_size=1, output_size=1)

# training
trainer = pl.Trainer(gpus=1, precision=16, max_epochs=100)
trainer.fit(model, train_loader)
```

PyTorch Lightning

PyTorch Lightning

Comparison PyTorch vs. PyTorch Lightning

Training

```
class LinearRegressionTorch(nn.Module):
    def __init__(self, input_size, output_size):
        super(LinearRegressionTorch, self).__init__()
        self.linear = nn.Linear(input_size, output_size)

    def forward(self, x):
        return self.linear(x)

model = LinearRegressionTorch(input_size=1, output_size=1)
model.train()

Zelle ausführen | Oben ausführen | Zelle debuggen
# %% Mean Squared Error
loss_fun = nn.MSELoss()

Zelle ausführen | Oben ausführen | Zelle debuggen
# %% Optimizer
learning_rate = 0.02
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

Zelle ausführen | Oben ausführen | Zelle debuggen
# %% perform training
losses = []
slope, bias = [], []
number_epochs = 1000
for epoch in range(number_epochs):
    for j, (X, y) in enumerate(train_loader):
        # optimization
        optimizer.zero_grad()
        # forward pass
        y_pred = model(X)
        # compute loss
        loss = loss_fun(y_pred, y)
        losses.append(loss.item())
        loss.backward()
        # update weights
        optimizer.step()
    # store loss
    losses.append(float(loss.data))
```

PyTorch

```
class LitLinearRegression(pl.LightningModule):
    def __init__(self, input_size, output_size):
        super(LitLinearRegression, self).__init__()
        self.linear = nn.Linear(input_size, output_size)
        self.loss_fun = nn.MSELoss()

    def forward(self, x):
        return self.linear(x)

    def configure_optimizers(self):
        learning_rate = 0.02
        optimizer = torch.optim.SGD(self.parameters(), lr=learning_rate)
        return optimizer

    def training_step(self, train_batch, batch_idx):
        X, y = train_batch

        # forward pass
        y_pred = model(X)

        # compute loss
        loss = self.loss_fun(y_pred, y)
        self.log('train_loss', loss, prog_bar=True)
        return loss

Zelle ausführen | Oben ausführen | Zelle debuggen
# %% model instance and training
# model instance
model = LitLinearRegression(input_size=1, output_size=1)

# training
trainer = pl.Trainer(gpus=1, precision=16, max_epochs=100)
trainer.fit(model, train_loader)
```

PyTorch Lightning

Early Stopping

Early Stopping

Additional Features: Early Stopping

```
from pytorch_lightning.callbacks.early_stopping import EarlyStopping
```

```
early_stop_callback = EarlyStopping(monitor="train_loss",  
min_delta=0.00, patience=2, verbose=True, mode="min")
```

```
trainer = pl.Trainer(accelerator='gpu', devices=1,  
max_epochs=500, log_every_n_steps=2, callbacks=  
[early_stop_callback])  
trainer.fit(model=model, train_dataloaders=train_loader)
```

